Recursive Self-Improvement for Robot Policies: Recursion Improves, the Gate Protects, the Meta Level Repairs

Anonymous CVPR submission

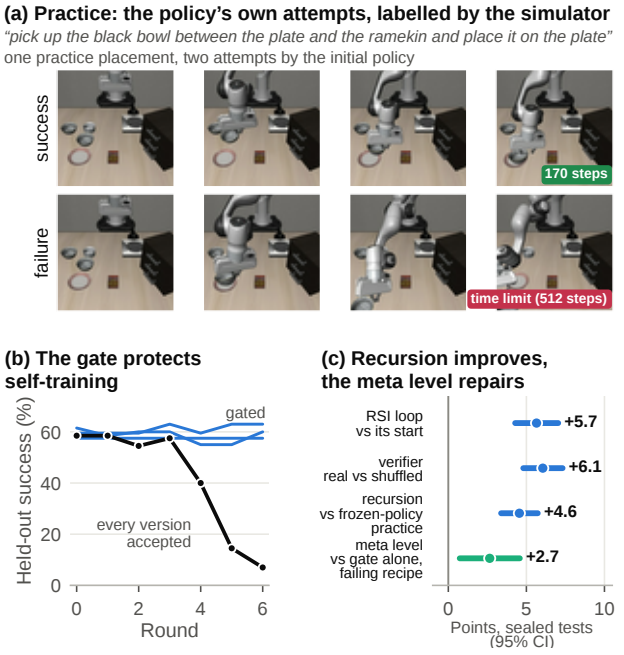


Figure 1. **Recursive self-improvement for a robot policy.** (a) One practice placement of a LIBERO-Spatial task and two attempts by the initial OpenVLA-OFT policy, as its camera saw them; the simulator's success check labels each attempt, and the loop trains on such labelled attempts. (b) LIBERO-Goal, behaviour cloning on the policy's own samples: when every new version is accepted, success on the round's 200 held-out placements falls from 58.5% to 7% within six rounds; under the gate the same recipe stays at its starting level. (c) On sealed tests: on 1,000 unseen LIBERO-Spatial scenes, the loop against its start, against the same loop with shuffled verifier labels, and against practice from the frozen initial policy; and the language-model meta level against the gate alone on the failing recipe of (b) (LIBERO-Goal, official test).

Abstract

Recursive self-improvement (RSI) lets a robot policy get better without new demonstrations: the policy practises on fresh scenes, a verifier labels its attempts, it trains on them, a gate decides whether the new version replaces the old one, and a meta level may rewrite the training procedure itself. We built an RSI framework for vision-language-action (VLA) policies in which every level is a switch, and measured each level's contribution with a 7B OpenVLA-OFT policy on two LIBERO suites under a preregistered protocol with two sealed tests. Recursion improves: *on 1,000 sealed unseen scenes the loop lifts the policy by* $+5.7$ *points* $[+4.3, +7.0]$ *on LIBERO-Spatial and* $+3.5$ $[+1.7, +5.4]$ *on LIBERO-Goal; without a real verifier, or without recursion, it ends 4.6 to 7.3 points lower.* The gate protects: *a self-training recipe that loses 15 points when every version is accepted holds its start under the gate,* $+13.6$ *and* $+13.0$ *points above plain training at equal compute on the two suites.* The meta level repairs: *from that recipe, a language model reading the loop's records diagnoses the fault in round one on every seed and lifts the loop* $+2.7$ *points* $[+0.7, +4.6]$ *over the gate alone, to the level of RL on the hand-tuned recipe without the pilot runs that tuned it. The limits are reported as registered: on the benchmark's fixed test states the gains are smaller (none on Spatial), the meta level adds nothing to a working recipe, and at equal training compute one round of plain RL on the tuned recipe ties the loop, or leads it by 2.7 points on Goal. We also find that a published VLA-RL score owes* $+16.8$ *points to training on its test states. Code, preregistrations and run records are released.*

1. Introduction

Collecting demonstrations is the expensive part of training a robot policy. An attractive alternative is to let the policy teach itself: give it a task, let it try, keep the attempts that worked, train on them, and let the improved policy try again. For language models this recipe works in several forms [3, 4, 6, 19, 21–23]. Robot learning has started to follow: a vision-language-action (VLA) policy can be improved by RL in simulation from a success flag alone [12, 14]. Recursive self-improvement goes one step further and also lets the system revise the procedure by which it improves.

Two questions are rarely answered. First, *which part does the work?* A self-improvement loop bundles a verifier

that labels attempts, a recursion in which the newest policy generates the next practice, a gate that keeps only versions that beat their parent, and sometimes a meta level that edits the recipe; it also contains ordinary RL as a special case, so a positive result does not show that any of these mattered. Second, *is the gain real?* LIBERO [13] tests on a fixed list of initial placements, and nothing stops a method from training on that list. We show that a widely used VLA-RL checkpoint does, and that about a fifth of its score comes from it (Sec. 5.5).

We answer both by experiment. Our framework (Sec. 3) has the three levels of an RSI system, and each is a switch: *recursion* (practice generated by the latest policy, or by the frozen initial one), *selection* (a gate on fresh paired placements, or none) and a *meta level* that may rewrite the training recipe (off, a numeric search, or a language model that reads the loop's records), plus a verifier that can be shuffled and a plain-RL baseline that spends the whole budget in one round. Every comparison, bar and statistic was committed publicly before its runs, the official test and 1,000 freshly sampled test scenes were sealed and read once, and every final averages three evaluations (Sec. 4).

What we found. Each level contributes something that training without it does not (Fig. 1, Tabs. 2 and 3).

1. **Recursion improves the policy.** On 1,000 sealed unseen scenes the loop is $+5.7$ points $[+4.3, +7.0]$ above its start on LIBERO-Spatial, on every seed and over its registered bar, and $+3.5$ $[+1.7, +5.4]$ on LIBERO-Goal after twelve rounds; with shuffled verifier labels, or with practice from the frozen initial policy, it ends 4.6 to 7.3 points lower (Sec. 5.1).
2. **The gate protects it.** Trained without the gate, a self-training recipe ends 14–16 points below its start on the official test, on every seed of both suites; under the gate it holds its start, $+13.6$ $[+10.5, +16.8]$ (Goal) and $+13.0$ $[+10.4, +15.6]$ (Spatial) above plain training at equal compute, and no gated run of 31 ends below its start beyond noise (Sec. 5.2).
3. **The meta level repairs the training procedure.** From the failing recipe, a language model reading the loop's records names the fault in round one on every seed and lifts the loop $+2.7$ $[+0.7, +4.6]$ over the gate alone, ending level with one-shot RL on the hand-tuned recipe ($+2.2$ $[+0.5, +3.9]$ on unseen scenes) without the $\geq$213 GPU-hours of pilots that tuned it (Sec. 5.3).
4. **The limits are real.** On the benchmark's fixed test states the gain over the start is smaller (zero on Spatial); the Goal improvement and the repair margin miss their bars ($+5$, $+3$); the meta level adds nothing to a working recipe; and one round of plain RL at equal compute ties the loop, and leads a strictly gated loop by 2.7 points on Goal's unseen scenes (Sec. 5.4).
5. **A published VLA-RL score owes about a fifth of itself to its test states** ($+16.8$ points; Sec. 5.5).

Contributions. (i) An RSI framework for VLA policies in which recursion, selection and a language-model meta level are switches inside one engine, with BC and GRPO as interchangeable updates. (ii) A preregistered protocol with two sealed tests and public amendments. (iii) The first controlled measurement of what each level of RSI contributes on a VLA benchmark, including where it falls short and why. (iv) A measurement of train-on-test inflation in a published VLA-RL result. Code, preregistrations and all run records are released.

2. Related work

Self-improvement with verifiers. STaR [22], ReST [6] and ReST-EM [20] fine-tune a language model on its own verified rationales; SPIN [3] and self-rewarding models [21] close the loop with the model as its own judge; Absolute Zero [23] lets the model write its own tasks and uses a code executor as the verifier. RLVR with group-relative policy optimisation [4, 19] is the dominant update rule. These works report that the loop helps; we ask which part of it helps. The iterated on-policy analysis of [8] predicts the early saturation of a learner restricted to samples from the initial policy, which our non-recursive control shows at twelve rounds.

Meta levels that edit the procedure. LLM-driven search over programs and training recipes has produced strong results when the search space is large and evaluation is cheap [7, 15, 18]. Recent outer-loop systems put the weight on the harness around the proposer: Meta-Harness [11] hands it the code, scores and raw traces of every earlier candidate as files instead of a summary, EvoTrainer [2] adds layered diagnostics and a memory that keeps failed branches, and LLMZero [5] has agents diagnose training signals at checkpoints. Our last meta level follows them (Sec. 3.4). In our setting evaluation is expensive and noisy and the space is a few dozen knobs; there, a meta level that diagnoses correctly still cannot beat a fixed recipe that works, and we isolate the case where it can.

RL for VLA policies. OpenVLA [9] and OpenVLA-OFT [10] are the policies we improve; SimpleVLA-RL [12], RL4VLA [14] and related work [16] apply on-policy RL to such policies in simulation with a binary success reward. We use the SimpleVLA-RL training stack as one update operator and its released checkpoints as starting points.

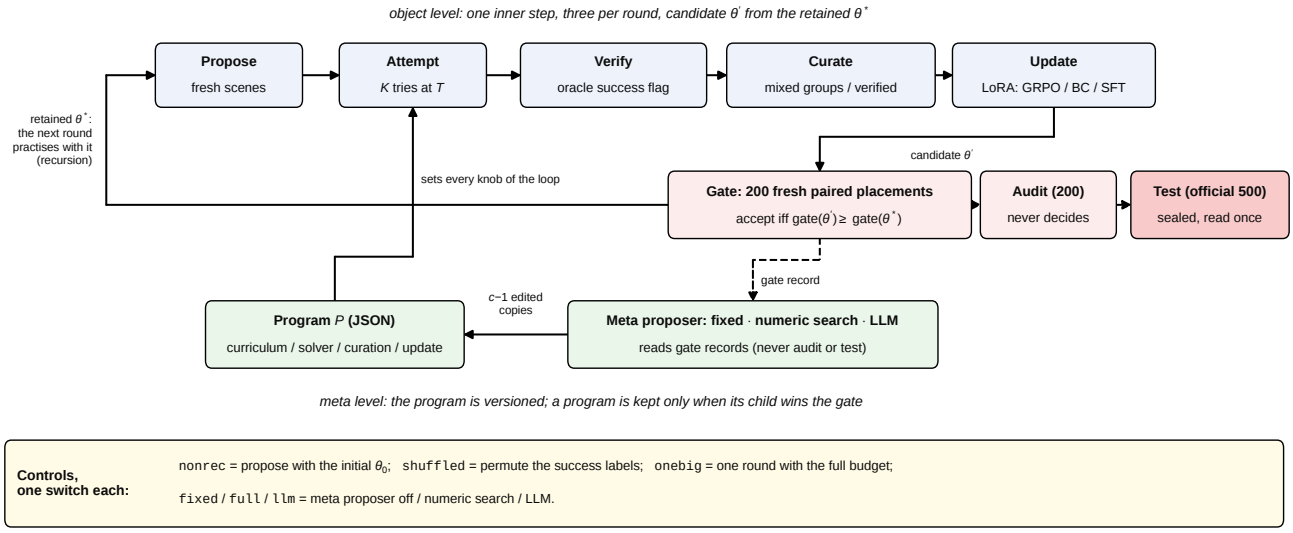


Figure 2. **The framework.** One engine, every ingredient a switch. The object loop practises on fresh scenes with the latest policy; the gate keeps a candidate only if it does not regress on paired held-out placements; the meta level versions the training program and may let a numeric search or an LLM edit it inside a declared space, seeing only gate records. Controls flip one switch each.

Preregistration. Preregistered designs are standard in the experimental sciences [17] and have been proposed for machine learning [1]; we follow the practice literally, with every change a dated public amendment.

3. Framework

Figure 2 shows the framework. A policy practises on freshly sampled scenes, a verifier marks each attempt, the verified attempts update the policy, and a gate decides whether the updated policy replaces the current one. An optional meta level above may edit the training recipe.

3.1. Setting and notation

A policy π_θ acts in scenes $s \sim \mathcal{S}$ (object placements of a task); an attempt $\tau \sim \pi_\theta(\cdot \mid s)$ is a trajectory, and an external oracle $v(s,\tau) \in \{0,1\}$ the policy cannot influence labels it (the simulator's success flag). A *program* P is a JSON document with four sections—`curriculum` (which scenes to practise, how many, how allocated across tasks), `solver` (attempts per scene K, temperature T), `curation` (which attempts become training items) and `update` (the operator and its hyperparameters)—and $\mathcal{P}$ is a declared edit space of allowed values. A *version* is a pair (θ, P). The framework keeps an archive of versions; the one it currently trusts is the *retained* version $(\theta^\star, P^\star)$.

3.2. The loop (Algorithm 1)

One *inner step* runs program P from policy θ: propose scenes, attempt each K times, verify, curate, update. Curation `mixed_groups` keeps every attempt of scenes whose success rate lies strictly inside a band (used by GRPO [19]: the group-centred advantage is zero when all attempts agree); `verified` keeps successful attempts only (behaviour cloning, BC). The update is a LoRA step on θ. A *round* is three inner steps from the same parent, producing a *candidate* θ'. Because the next round starts from the retained version, practice is always generated by the latest policy: this is the *data recursion*, and the control `nonrec` breaks exactly it by proposing and attempting with the frozen initial policy θ_0 while still updating from $\theta^\star$.

3.3. Gate, audit and sealed test

Three disjoint evaluation sets are drawn per run from seeds that never overlap the official benchmark states. The *gate* (20 placements per task, 200 episodes, greedy decoding, environment re-seeded per episode so that versions are compared on identical placements) decides: a candidate is kept iff its gate success is at least the parent's on the same placements. In the *fresh-paired* mode used in every main run, a new gate set is drawn every round and the parent is re-evaluated on it, so acceptance is always a paired comparison on placements never seen before and a version cannot overfit a fixed gate. Because the gate's verdict decides the next version, and that version generates the next practice, the loop is recursive in selection as well as in data. The *audit* (200 more episodes) is recorded but never decides anything; it exposes gate overfitting. The *test* is the official benchmark (50 initial states per task, 500 episodes) and is *sealed*: no process in the loop can read it, and it is opened once, after every decision in every run of a study has been made. Be-

Algorithm 1 Gated recursive self-improvement with a meta level

Require: initial policy θ_0, seed program P_0, oracle v, gate set size n_g, rounds R, candidates per round c

$(\theta^\star, P^\star) \leftarrow (\theta_0, P_0)$; evaluate $\theta^\star$ on fresh gate and audit sets
for $r = 1 \dots R$ **do**
  draw a fresh gate set G_r of n_g paired placements; re-evaluate $\theta^\star$ on G_r
  $\{P_1, \dots, P_c\} \leftarrow \{P^\star\} \cup$ PROPOSE(round records, $\mathcal{P}$) ▷ meta level: fixed / numeric / LLM
  for each candidate program P_i **do**
    $\theta_i \leftarrow \theta^\star$
    for inner step $= 1 \dots 3$ **do**
      scenes $\leftarrow$ PROPOSE$(P_i, \theta_{\text{gen}})$ with $\theta_{\text{gen}} = \theta_i$ (recursive) or θ_0 (`nonrec`)
      attempts $\leftarrow K$ samples per scene at temperature T; labels $\leftarrow v(\cdot)$ (or permuted: `shuffled`)
      items $\leftarrow$ CURATE(attempts, P_i); $\theta_i \leftarrow$ UPDATE(θ_i, items, P_i)
    end for
    $g_i \leftarrow$ success of θ_i on G_r (greedy); record audit, never use it
  end for
  $i^\star \leftarrow \arg\max_i g_i$; **if** $g_{i^\star} \geq g(\theta^\star, G_r)$ **then** $(\theta^\star, P^\star) \leftarrow (\theta_{i^\star}, P_{i^\star})$ ▷ gate
end for
read the sealed test of $\theta^\star$ once, after every run of the study has ended

cause the loop practises on sampled placements while the official states are a fixed list, a second sealed test, registered before any of its evaluations, scores each retained policy on 1,000 freshly sampled placements that no run trained, gated or audited on.

3.4. Meta level

The meta level versions the program and lets it be improved under the same gate. Each round the incumbent $P^\star$ and $c-1$ proposed programs are each trained for real from the same parent, and the gate picks the version to keep; a program is credited only when its child wins. Three proposers:

- `fixed`: no change (the meta level is off).
- `numeric`: random mutation of declared knobs, biased toward knobs whose past changes raised the gate.
- `llm`: a language model (Claude Sonnet) reads the current program, the edit space with a neutral description of what each option computes, and, for every past candidate, its program diff, gate score, acceptance and per-step statistics (sampled success, share of mixed scenes, item counts, first/last training loss). It must write a diagnosis, then return one new program.
- `llm+harness`: the same model inside a harness, following [2, 11]: a *workspace* of files (the full record of every candidate of the run, per-task gate tables, an evidence ledger that pools gate changes per edited field), a frozen *memory* of earlier runs with a playbook the model distilled before launch, its own notes from earlier rounds, and list, read and search tools used over several turns. The prompt states the measured gate noise. Building it exposed a defect in the plain `llm` proposer: the training loss it was told to watch never reached it (a field-name error), in every run reported for it below.

Table 1. **Controls.** Each arm flips one switch of Algorithm 1.

arm	change in Algorithm 1	isolates
`fixed`	none (proposer off)	gated recursive loop
`nonrec`	$\theta_{\text{gen}} = \theta_0$	data recursion
`shuffled`	labels permuted per scene	the verifier
`onebig`	R=1, $3R$ inner steps	gate + rounds vs one-shot RL
`full`	numeric proposer, c=2	meta level (equal compute)
`llm`∗	LLM proposer, c=2	meta level (equal compute)

∗ `llm`: knob space; `llm_harness`: harness space; `llm_harness_cap`: harness space under the compute cap.

The *knob* space has ten hyperparameters (scenes, allocation, samples, temperature, curation rule and budget, operator, learning rate, steps, batch); the *harness* space adds nine components of the procedure (policy-gradient ratio, loss aggregation, BC likelihood temperature, advantage normalisation, curation band, per-scene cap, allocation floor, two clip bounds). Each corresponds to a change made by hand during development after diagnosing round records; the LLM is never told which option is better, so rediscovering them is a test of the meta level.

Attribution guard. The meta level never sees the audit, the test, the weights or the episodes; it only chooses values inside $\mathcal{P}$, every proposal passes the same validator as a numeric mutation, failures fall back to numeric search and are counted, and an optional *compute cap* keeps any candidate within the starting program's rollout and optimiser budget. All practice data comes from the policy's own attempts and all verdicts from the oracle; the LLM chooses procedure, not data.

3.5. Every control is a switch

Table 1 lists the arms used in the experiments. Each differs from `fixed` in exactly one line of Algorithm 1, so a difference between an arm and `fixed` is attributable to that line; `full` and the `llm` arms train the same number of candidates and so cost the same.

Table 2. **What the loop and its ingredients add, on two suites and two sealed tests** (points, 95% CI, seed-averaged paired bootstrap; bold: registered bar met). Official: the benchmark's 500 fixed states, three-evaluation finals against the pooled initial reference. Fresh: 1,000 sampled placements, disjoint from every practice, gate and audit set and registered before any evaluation on them. Goal: twelve rounds, seeds 1–3; Spatial: six rounds, seeds 0–3 (start vs. official reference: seeds 1–3).

		LIBERO-Goal, 12 rounds		LIBERO-Spatial, 6 rounds	
comparison	bar	official	fresh scenes	official	fresh scenes
loop − start	+5 / +3†	+1.60 [−0.88, +4.13]	+3.54 [+1.72, +5.42]	−0.01 [−2.15, +2.13]	**+5.65 [+4.28, +7.04]**
recursion: loop − `nonrec`	+2	**+7.84 [+5.22, +10.49]**	**+7.30 [+5.47, +9.20]**	**+5.73 [+4.12, +7.38]**	**+4.55 [+3.38, +5.73]**
verifier: loop − `shuffled`	+3	–	–	+1.67 [+0.18, +3.13]	**+6.05 [+4.80, +7.32]**
vs. plain RL: loop − `onebig`	+2	−0.38 [−2.42, +1.64]	−2.73 [−4.30, −1.20]	−0.45 [−1.83, +0.93]	+0.38 [−0.60, +1.38]

† +5 on the official test and on Goal's fresh scenes; +3 on Spatial's fresh scenes (registered before any Spatial run ended), where +5.65 also clears +5.

Table 3. **Registered claims on LIBERO-Goal** (official 500-episode test, 3-evaluation finals, pooled initial reference, seed-averaged paired bootstrap). Platform Y = cloud pool, A = A800.

claim	comparison	points	95% CI	bar	seeds
R1	`fixed` − ref. (Y)	+1.71	[−0.76, +4.23]	+5	1–3
R1	`full` − ref. (Y)	+1.28	[−1.37, +3.99]	+5	1–3
R2	`fixed` − `shuffled` (Y)	+1.22	[−0.37, +2.85]	+3	0–3
R3	`fixed` − `nonrec` (Y)	+1.72	[−0.28, +3.68]	+2	0–3
R4	`full` − `fixed` (Y)	−0.42	[−2.40, +1.56]	+2	1–3
R4	`llm_harness` − `fixed` (Y)	−0.73	[−2.58, +1.07]	+2	1–3
R4	`llm` − `fixed` (A)	+1.04	[−0.84, +3.00]	+2	1–3
R4	`llm_harness_cap` − `fixed` (A)	+1.16	[−0.76, +3.16]	+2	1–3
R5	best gated − `onebig` (Y)	−0.58	[−2.12, +0.95]	+2	0–3
noise	half vs. half of the n0 evaluations (Y / A)	+0.64 / +0.54	[−0.07, +1.36] / [−0.24, +1.35]		

4. Experimental setup

4.1. Benchmarks, policies and starting points

LIBERO / OpenVLA-OFT. LIBERO-Goal [13] has ten goal-conditioned tasks in one scene, LIBERO-Spatial ten pick-and-place tasks that differ in the spatial layout of the same objects, and LIBERO-10 ten long-horizon tasks. The policy is OpenVLA-OFT [10] (7B, action chunks of 8), started from a supervised fine-tune on *one* demonstration per task [12] so that there is room to improve (about 59% on LIBERO-Goal, 60% on LIBERO-Spatial, 19% on LIBERO-10). Rollouts and GRPO/BC training use the SimpleVLA-RL stack with LoRA (r=32). The seed program—chosen from sealed pilots and frozen by amendment before any main run—is GRPO with learning rate 3×10^{-5}, matched-ratio log-probabilities, episode-balanced loss, 8 samples per scene at temperature 1.5, 160 scenes per inner step, 200 optimiser steps of batch 64; 6 rounds of 3 inner steps; fresh paired gates.

4.2. Arms and baselines

The LIBERO-Goal main study runs the eight arms of Sec. 3.5 with three seeds each (a fourth seed was added to the `nonrec`, `shuffled`, `onebig` and `fixed` arms by amendment before any of their runs ended), 31 runs in all, on two GPU platforms: a shared 8×A800 node and a 16-GPU cloud pool with its own evaluation stack. The stacks are not bit-identical, so every comparison is paired within a platform and the platforms serve as replicates. `onebig` is the plain-RL baseline: the whole six-round rollout and update budget in one round of 18 on-policy inner steps from the initial policy, with the same sampling and update code; it is read both through the gate and as its ungated single candidate.

A second study on LIBERO-Spatial, registered before any of its runs, repeats the four object-level arms (`fixed`, `nonrec`, `shuffled`, `onebig`) with four seeds under the Goal bars (sixteen runs); a twelve-round extension on Goal (`fixed12`, `nonrec12`, `onebig12`) and a self-repair study (Sec. 5.3) complete the design.

4.3. Preregistered protocol

We fixed the protocol before the runs and release the pre-registration documents and every dated amendment with the code.

Sealing. The official LIBERO test states are never used for training, sampling, gating or program selection, and no test file is opened before a study's read-once. Only two

sealed pilots per model were allowed before the seed program was frozen.

Finals and reference. The *same* initial policy scored between 57.2% and 60.0% in different runs, and about 40% of episodes flip between nearby policies. Each run's final is therefore the per-episode mean of three evaluations of its retained policy, and the initial reference is the per-episode mean of every in-run evaluation of the initial policy on that platform (36 on the cloud pool, 21 on the A800). A noise check (the same statistic between two halves of the initial-policy evaluations) accompanies every verdict.

Claims and bars. Seed-averaged per-episode paired bootstrap (10,000 draws), 95% CI. "CI > 0" means the interval excludes zero.

claim	comparison	bar (points)
R1 improvement	`fixed`, `full` − reference	$\geq +5$, CI > 0, 2/3 seeds > 0
R2 verifier	`fixed` − `shuffled`	$\geq +3$, CI > 0
R3 recursion	`fixed` − `nonrec`	$\geq +2$, CI > 0
R4 meta level	each meta arm − `fixed`	$\geq +2$, CI > 0
R5 vs one-shot RL	best gated arm − `onebig`	$\geq +2$, CI > 0
M2-A meta v2	`llm` − identical-recipe replica	$\geq +3$, CI > 0
S1 self-repair	`llm_harness` − `fixed` (failing start)	$\geq +3$, CI > 0

Negative readings are reported as such; the reading scripts were committed before each read and their outputs are released.

5. Results

We report what each level of the architecture contributes (Fig. 2): recursive practice (Sec. 5.1), the gate (Sec. 5.2) and the meta level (Sec. 5.3), then how the loop compares with plain RL (Sec. 5.4). Every number comes from a pre-registered read of a sealed test unless it is marked *gate-level*.

5.1. Does recursive self-improvement make the policy better?

Finding 1: on scenes it never saw, the loop improves the policy on both suites, and recursion and the verifier are what make it work. Table 2 collects the registered reads. On LIBERO-Spatial (sixteen runs) the loop lifts the policy by $+5.65$ points $[+4.28, +7.04]$ on 1,000 sealed fresh scenes, on every seed ($+4.8$ to $+6.8$), clearing its registered bar. Each defining ingredient is needed: with shuffled verifier labels the loop ends 6.05 points lower, at its start, and with practice generated by the frozen initial policy it ends 4.55 points lower; recursion also passes on the official test ($+5.73$ $[+4.12, +7.38]$), already at six rounds. On LIBERO-Goal the same pattern needs twelve rounds: the loop ends $+3.5$ points $[+1.7, +5.4]$ above its start on the fresh scenes (short of its $+5$ bar), and recursion is worth $+7.3$ there and $+7.8$ on the official test, positive on every seed. Without recursion the Goal loop falls *below* its start (-3.8 on the fresh scenes) although every version passed a fresh paired gate: practice generated by the initial policy no longer matches the learner, its candidates are truly a little worse than their parent, and noise-driven acceptances accumulate. Recursive practice keeps the data on the learner's own distribution. On LIBERO-Object one always-accept GRPO curve (one seed, held-out placements) rises from 55.8% to 82.0%.

Finding 2: the official fixed states show less. On Spatial's 500 official states the same policies show no gain over the reference (-0.01 $[-2.15, +2.13]$) while recursion and the verifier keep their sign ($+5.73$, $+1.67$ $[+0.18, +3.13]$), and on Goal at six rounds the registered gains are one to two points, none over its bar (Tab. 3; R1 $+1.71$ $[-0.76, +4.23]$ against $+5$). At the gate (Fig. 3) the Goal arms separate within six rounds.

5.2. Is self-training safe? The gate

Finding 3: the gate is what keeps self-training from degrading. Trained without the gate, the behaviour-cloning recipe of Sec. 5.3 ends 15.2 points $[11.5, 18.8]$ below its start on the official test and 16.3 below on the fresh scenes, on all three seeds (one earlier always-accept run fell to 5.5%, Fig. 1b). At equal training compute the gated loop ends $+13.6$ $[+10.5, +16.8]$ and $+15.5$ $[+13.2, +17.8]$ above it, past the registered $+10$ bars, because it holds its start. On Spatial the same comparison gives $+13.0$ $[+10.4, +15.6]$ on the official test (bar met) and $+8.4$ $[+6.7, +10.1]$ on the fresh scenes (under it); plain training loses 15.0 and 7.2 points there (1.2% on one A800 seed). The gate is cheap (one extra evaluation per candidate) but noisy (sd about 3.4 points at 200 episodes), and its protection relies on recursive practice (Finding 1).

5.3. Can the system improve its own training procedure? The meta level

A recipe that fails. The meta level can be told apart from RL only where RL's recipe fails. One such case was found before any meta-level run: behaviour cloning on the policy's own temperature-1.6 samples with a unit-temperature likelihood degrades when its versions are accepted unconditionally. The repair is a one-knob change (the BC temperature, the sampling temperature or the update operator), and all three knobs are in the harness space, but the meta level is never told which. A registered study starts from this recipe with arms `fixed` (the gate alone), `full` (numeric search) and `llm_harness` (same compute), four seeds, and the bar S1 $=$ `llm_harness` $-$ `fixed` $\geq +3$, CI above 0.

Finding 4: the meta level diagnoses and repairs a failing training procedure on its own. On all four seeds the language model's first-round diagnosis names the mismatch between the sampling and likelihood temperatures and the missing failure signal (verbatim in the released logs), and its first proposal is the same: switch to GRPO on mixed-

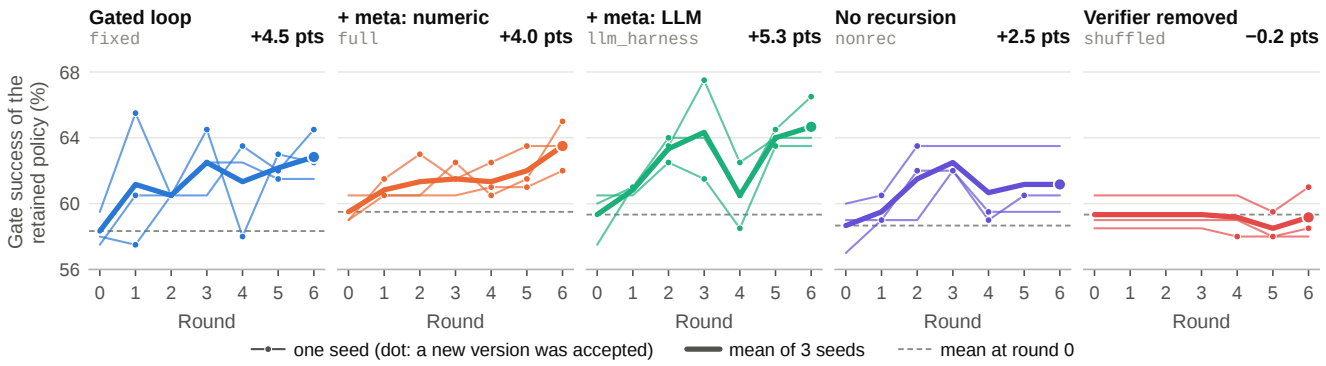


Figure 3. **Gate success of the retained policy by round** (LIBERO-Goal, cloud platform; thin lines are seeds 1–3, the thick line their mean). A dot marks a round that accepted a new version; every round is scored on 200 newly drawn placements, so a retained version's score can fall from one round to the next. The number in each panel is the change of the mean from round 0 to round 6. With a real verifier and recursive practice the loop climbs 4–5 points at the gate within six rounds; without the verifier (`shuffled`) it stays flat, and with practice from the frozen initial policy (`nonrec`) it stops after round 3. Gate gains overstate test gains by the winner's curse.

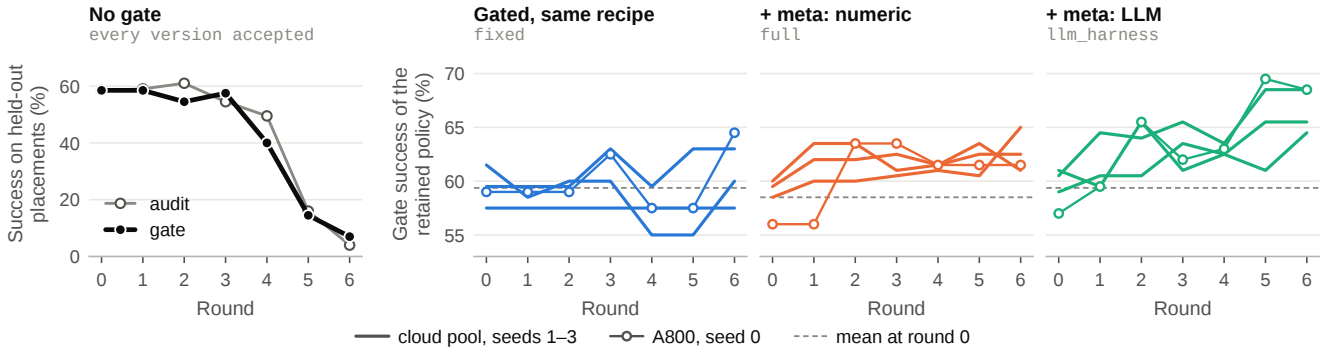


Figure 4. **Self-repair study, gate level.** Left: the starting BC recipe collapses when every version is accepted. Right (narrower scale): under the gate the same recipe stays at its starting level (`fixed`); both proposers find repairs; at the gate the LLM arm looks 4–8 points ahead, on the sealed test it is +2.65 (Tab. 4).

Table 4. **Self-repair study**: gate / sealed-test final (%, three evaluations) of each retained version, and what changed. S1 = `llm_harness` − `fixed` = +2.65 [+0.73, +4.58] (bar +3 missed); S2 = `llm_harness` − `full` = +1.15 [−0.63, +2.92]; S3 = `llm_harness` − hand-tuned GRPO loop = +0.28 [−1.43, +1.97].

seed	`fixed` (gate only)	`full` (numeric)	`llm_harness`
1 (Y)	60.0 / 56.1	62.5 / 56.7	65.5 / 60.1, GRPO in r1
2 (Y)	start / 58.7, all refused	65.0 / 59.9, GRPO	68.5 / 61.3, GRPO in r1
3 (Y)	63.0 / 57.5	61.0 / 60.2, BC rollout-T	64.5 / 57.9, BC T 1.6→1.3
0 (A)	64.5 / 57.4	61.5 / 58.8	68.5 / 60.9, GRPO in r2

outcome scenes at a lower temperature. The gate kept that switch in round 1–2 on three seeds and refused it twice on the fourth, where the model repaired BC instead (Tab. 4). The repaired loop then gains what the hand-tuned loop gains: S1 = +2.65 points [+0.73, +4.58] over the gate alone, positive on three seeds with the interval above zero, and level with the hand-tuned GRPO loop (S3 = +0.28 [−1.43, +1.97]), while the gate alone holds the policy at its start (−2.37 [−4.12, −0.62] against the hand-tuned loop). S1 is under its +3 bar, so the claim is not made in its registered form, and numeric search over the same knobs also finds a repair, on two seeds and only in round 6, ending 1.15 points behind ([−0.63, +2.92]): the model's advantage is the speed and reliability of its diagnosis, not a better end point. Six gate-only pilots found no second failing recipe.

Finding 5: on a recipe that already works, the meta level finds nothing to repair. All four meta-level arms of the main study land within ±1.2 points of `fixed` (R4, Tab. 3). A second registered study gives the model per-task gate rates and compares it with a *replica* arm training the identical recipe: −0.27 [−2.53, +2.00] and +0.30 [−1.60, +2.15] on the two platforms. The model's diagnoses are correct and its edits are those a practitioner would try, but its proposals scored −1.2 points against the incum-

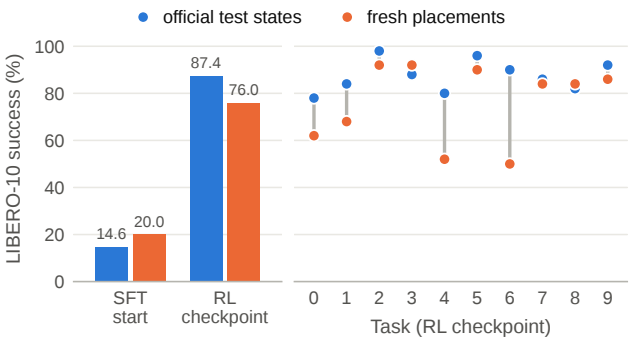


Figure 5. **Train-on-test inflation** (LIBERO-10). Left: official vs. fresh placements for the SFT start and the RL checkpoint trained on the official states. Right: the RL checkpoint per task; the loss is concentrated on tasks 4 and 6.

bent at the gate (n=30) where replicas scored $+0.6$ (sd 3.4, n=36): *every recipe in the space reaches the same plateau*, so there is nothing for a correct diagnosis to fix.

5.4. Against plain RL at equal compute

Finding 6: a strictly gated loop ties plain RL on a recipe that works, and can trail it. One round of plain RL with the whole budget (`onebig`) ties the loop on Spatial on both tests (-0.45, $+0.38$) and on Goal at six rounds (R5 -0.58), ties the twelve-round Goal loop on the official test (-0.38) and is 2.7 points ahead of it on the fresh scenes (Tab. 2). The cost is the gate's strictness: the twelve-round runs kept 6, 9 and 9 of 12 rounds, mostly refusing dips within the gate's noise. Counting the tuning changes the picture. The recipe both sides use was chosen by 21 pilot runs ($\geq$213 GPU-hours, about 13 one-shot runs); started from the failing recipe instead, the meta level ends level with one-shot RL on the tuned recipe on the official test ($+0.05$ $[-1.6, +1.7]$) and $+2.2$ $[+0.5, +3.9]$ ahead on the fresh scenes, with twice one run's training and no pilots (descriptive).

5.5. How much of a published VLA-RL number is the test set?

Finding 7: about a fifth of the headline number. The released SimpleVLA-RL LIBERO-10 checkpoint was trained by RL on the 50 fixed test states per task. On the 500 official episodes and on 500 placements freshly drawn by LIBERO's own sampler, its SFT start scores 14.6% and 20.0%, the RL checkpoint 87.4% and 76.0%. The difference of the two gaps, which cancels any difference in difficulty between the placement sets, is $+16.8$ points $[+9.8, +23.4]$; 500 further never-seen placements give $+16.0$ $[+9.2, +22.4]$ (Fig. 5). The method works ($+56$ points on fresh placements), but about a fifth of its headline number is memorisation of the test states.

6. Discussion

What the evidence does not support. The improvement over the start clears its bar only on Spatial's unseen scenes, not on the official test of either suite nor on Goal's unseen scenes, and the self-repair margin misses its bar ($+2.7$ against $+3$). We do not claim that the loop beats tuned RL at equal training compute (it reaches tuned RL without the tuning), nor that the language model's diagnosis, rather than the knob space it edits, carries the repair (its margin over blind search is $+1.2$ $[-0.6, +2.9]$).

For practitioners. A fresh paired gate makes an unsafe recipe safe for one evaluation per candidate if its strictness matches its noise; fresh test scenes should complement fixed test states.

Limitations. One simulated benchmark family, one 7B policy and a few-demonstration start; other policies and real robots are untested. Preregistration constrains the analysis, not the design: a different recipe space might leave more for a meta level to find, and "without tuning" means without pilot runs, not without our choice of that space.

7. Conclusion

Each level of recursive self-improvement does a distinct job for a robot policy: recursion makes it better on unseen scenes of two suites, the gate keeps self-training from degrading, and the meta level repairs a failing training procedure. We hope its protocol of sealed tests and public amendments outlives the numbers.

References

[1] Samuel Albanie, James Thewlis, Sébastien Ehrhardt, and João F. Henriques. On the preregistration of machine learning experiments. *NeurIPS Workshop on Preregistration in ML*, 2021. 3

[2] Guhong Chen, Yingcheng Shi, Yongbin Li, Binhua Li, Xander Xu, Hu Wei, Shiwen Ni, Min Yang, and Jieping Ye. EvoTrainer: Co-evolving LLM policies and training harnesses for autonomous agentic reinforcement learning. *arXiv:2606.03108*, 2026. 2, 4

[3] Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-play fine-tuning converts weak language models to strong language models. In *ICML*, 2024. 1, 2

[4] DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv:2501.12948*, 2025. 1, 2

[5] Haoyang Fang, Wei Zhu, Boran Han, Alex Zhang, Zhenyu Pan, Shuo Yang, Shuai Zhang, Jiading Gai, Peng Tang, Cuixiong Hu, Xuan Zhu, Huzefa Rangwala, George Karypis, and Bernie Wang. LLMZero: Discovering adaptive training strategies for RL post-training via LLM agents. *arXiv:2606.18388*, 2026. 2

[6] Caglar Gulcehre, Tom Le Paine, Srivatsan Srinivasan, Ksenia Konyushkova, Lotte Weerts, Abhishek Sharma, Aditya Siddhant, Alex Ahern, Miaosen Wang, Chenjie Gu, et al. Reinforced self-training (ReST) for language modeling. *arXiv:2308.08998*, 2023. 1, 2

[7] Shengran Hu, Cong Lu, and Jeff Clune. Automated design of agentic systems. In *ICLR*, 2025. 2

[8] Juno Kim, Jihun Yun, Jason D. Lee, and Kwang-Sung Jun. Coverage improvement and fast convergence of on-policy preference learning. *arXiv:2601.08421*, 2026. 2

[9] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, et al. OpenVLA: An open-source vision-language-action model. In *CoRL*, 2024. 2

[10] Moo Jin Kim, Chelsea Finn, and Percy Liang. Fine-tuning vision-language-action models: Optimizing speed and success. *arXiv:2502.19645*, 2025. 2, 5

[11] Yoonho Lee, Roshen Nair, Qizheng Zhang, Kangwook Lee, Omar Khattab, and Chelsea Finn. Meta-Harness: End-to-end optimization of model harnesses. *arXiv:2603.28052*, 2026. 2, 4

[12] Haozhan Li, Yuxin Zuo, Jiale Yu, Yuhao Zhang, Zhaoyang Yang, Kaiyan Zhang, Xuekai Zhu, Yuchen Zhang, Tianxing Chen, Ganqu Cui, et al. SimpleVLA-RL: Scaling VLA training via reinforcement learning. *arXiv:2509.09674*, 2025. 1, 2, 5

[13] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. In *NeurIPS Datasets and Benchmarks*, 2023. 2, 5

[14] Jijia Liu, Feng Gao, Bingwen Wei, Xinlei Chen, Qingmin Liao, Yi Wu, Chao Yu, and Yu Wang. What can RL bring to VLA generalization? an empirical study. *arXiv:2505.19789*, 2025. 1, 2

[15] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI scientist: Towards fully automated open-ended scientific discovery. *arXiv:2408.06292*, 2024. 2

[16] Guanxing Lu, Wenkai Guo, Chaoyi Zhang, Yue Zhou, Haonan Jiang, Zifeng Gao, Yansong Tang, and Ziwei Wang. VLA-RL: Towards masterful and general robotic manipulation with scalable reinforcement learning. *arXiv:2505.18719*, 2025. 2

[17] Brian A. Nosek, Charles R. Ebersole, Alexander C. DeHaven, and David T. Mellor. The preregistration revolution. *PNAS*, 115(11):2600–2606, 2018. 3

[18] Alexander Novikov, Ngân Vũ, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, Sergey Shirobokov, Borislav Kozlovskii, Francisco J. R. Ruiz, Abbas Mehrabian, et al. AlphaEvolve: A coding agent for scientific and algorithmic discovery. *arXiv:2506.13131*, 2025. 2

[19] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv:2402.03300*, 2024. 1, 2, 3

[20] Avi Singh, John D. Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, Peter J. Liu, James Harrison, Jaehoon Lee, Kelvin Xu, et al. Beyond human data: Scaling self-training for problem-solving with language models. In *TMLR*, 2024. 2

[21] Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Xian Li, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. Self-rewarding language models. In *ICML*, 2024. 1, 2

[22] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Bootstrapping reasoning with reasoning. In *NeurIPS*, 2022. 2

[23] Andrew Zhao, Yiran Wu, Yang Yue, Tong Wu, Quentin Xu, Matthieu Lin, Shenzhi Wang, Qingyun Wu, Zilong Zheng, and Gao Huang. Absolute zero: Reinforced self-play reasoning with zero data. *arXiv:2505.03335*, 2025. 1, 2